

CONTENTS INCLUDE:

- Distributed Version Control Concepts
- Cloning Existing Projects
- The Typical Local Workflow
- The Remote Workflow
- CVS & Subversion Interoperability
- and more...

Getting Started with **Git**

By Matthew McCullough

WHY GET GIT?

Git is a postmodern version control system that offers the familiar capabilities of CVS or Subversion, but doesn't stop at just matching existing tools. Git stretches the very notion of version control systems (VCS) by its ability to offer almost all of its features for use offline and without a central server. It is the brainchild of Linus Torvalds, with the first prototype written in a vitriolic two-week response to the "BitKeeper debacle" of 2005.

Today, developers everywhere are migrating in droves to this exciting platform. Users reference its blistering performance, usage flexibility, offline capabilities, and collaboration features as their motivation for switching. Let's get started with Git. You'll be using it like a master in no time at all.


More Than Just a Version Control System

Though originally targeted at Linux kernel developers, Git has found a home in many other version-needing systems, such as document revision control. Git is just as capable of versioning a folder of documents or configuration files as it is of tracking decades of source code for the Perl project. This has led to Git being used by writers and network administrators — a case less likely to be encountered with the cumbersome CVS and Subversion servers and their ceremonious setup.

DISTRIBUTED VERSION CONTROL

If you are familiar with one or more traditional or centralized version control systems like Subversion, there will be several mental adjustments to make in your migration to Git. The first is that there is no central server. The second is that there is no central server. The **full history** of the repository lives on every user's machine that has cloned (checked out) a copy of the repository. This is the essence of a Distributed Version Control System (DVCS).

Once over those hurdles, it is quite liberating to be able to work entirely independently, versioning any new project that you start, even if in the incubation phase. The ease of setting up a new Git repository (or 'repo' in common parlance) leads to setting up repos everywhere. It feels frictionless.



From there you'll progress to the second epiphany of being able to share a repository and a changeset directly with a colleague without any complicated setup, without a checkin to a central server, direct network connectivity, or having to worry about firewalls getting in the way. Git has done technologically for version control what BitTorrent did for file sharing. It permanently replaced the spoke and hub structure with a peer-to-peer model, and there's no turning back. It supports transmitting binary sets of changes via USB stick, email, or in the traditional style, over a network, but amazingly, via HTTP, FTP, SCP, Samba, SSH, or WebDAV.

GETTING STARTED

Installing Git

Git has a very light footprint for its installation. For most platforms, you can simply copy the binaries to a folder that is on the executable search \$PATH. Git is primarily written in C, which means there is a unique distribution for each supported platform.



The canonical reference for Git installers can be found on a subpage of the official Git site. <http://git-scm.com/download>

Establishing User Credentials

Once you have selected a suitable distribution of Git for your platform, you'll need to identify yourself with a username and email address to Git.

In a separation of concerns most satisfying to the purist, Git does not directly support repository authentication or authorization. It delegates this in a very functional way to the protocol (commonly SSH) or operating system (file system permissions) hosting or serving up the repository. Thus, the

Don't Miss An Issue!

Get over 90 DZone Refcardz
FREE from Refcardz.com!



New Release
Every Monday

Visit Refcardz.com to get them all Free!

user information provided during your first Git setup on a given machine is purely for “credit” of your code contributions.

With the binaries on your \$PATH, issue the following three commands just once per new machine on which you’ll be using Git. Replace the username and email address with your preferred credentials.

```
git config --global user.name "matthew.mccullough"
git config --global user.email "matthewm@ambientideas.com"
git config --global color.ui "auto"
```

These commands store your preferences in a file named **.gitconfig** inside your home directory (~ on UNIX and Mac, and %USERPROFILE% on Windows).

If you are intrigued by all the potential nuances of a Git setup, GitHub, a web-based code hosting site, offers several in-depth tutorials on setting up Git for Linux, Windows, and Mac. Here are several in-depth Git installation guides:

```
http://help.github.com/win-git-installation/
http://help.github.com/mac-git-installation/
http://help.github.com/linux-git-installation/
```

Creating a Repository

Now that Git is installed and the user information established, you can begin establishing new repositories. From a command prompt, change directories to either a blank folder or an existing project that you want to put under version control. Then initialize the directory as a Git repository by typing the following commands:

```
git init
git add .
git commit -m 'The first commit'
```

The first command in the sequence, **init**, builds a **.git** directory that contains all the metadata and repository history. Unlike many other version control systems, Git uniquely stores everything in just a single directory at the top of the project. No pollution in every directory.

Next, the **add** command with the dot wildcard tells Git to start tracking changes for the current directory, its files, and for all folders beneath, if any exist.

Lastly, the **commit** function takes all previous additions and makes them permanent in the repository’s history in a transactional action. Rather than letting Git prompt the user via the default text editor, the **-m** option preemptively supplies the commit message to be saved alongside the committed files.

It is amazing and exciting to be able to truthfully say that you can use the basics of Git for locally versioning files with just these three commands.

CLONING EXISTING PROJECTS

An equally common use case for Git is starting from someone else’s repository history. This is similar to the checkout concept in Subversion or other centralized version control systems. The difference in a DVCS is that the **entire history**, not just the latest version, is retrieved and saved to the local user’s disk.

The syntax to pull down a local copy of an existing repo is:

```
git clone git://github.com/matthewmccullough/hellogitworld.git
or
git clone http://github.com/matthewmccullough/hellogitworld.git
or
git clone git@github.com:matthewmccullough/hellogitworld.git
```

The protocol difference often signifies whether you have read-only or writeable access to the origin repository. The final syntax, which accesses an SSH exposed repository, is the most common write-enabled protocol.

The **clone** command performs several subtasks under the hood. It sets up a **remote** (a Git repository address bookmark) named origin that points to the location

git://github.com/matthewmccullough/hellogitworld.git. Next, **clone** asks this location for the contents of its entire repository. Git copies those objects in a zlib-compressed manner over the network to the requestor’s local disk. Lastly, **clone** switches to a branch named **master**, which is equivalent to Subversion’s **trunk**, as the current working copy. The local copy of this repo is now ready to have edits made, branches created, and commits issued – all while online or offline.

TREEISH & HASHES

Rather than a sequential revision ID, Git marks each commit with a SHA-1 hash that is unique to the person committing the changes, the folders, and the files comprising the changeset. This allows commits to be made independent of any central coordinating server.

A full SHA-1 hash is 40 hex characters
64de179becc3ed324daab72f7238df1404723672

To efficiently navigate the history of hashes, several symbolic shorthand notations can be used as listed in the table below. Additionally, any unique sub-portion of the hash can be used. Git will let you know when the characters supplied are not enough to be unique. In most cases, 4-5 characters are sufficient.

Treeish	Definition
HEAD	The current committed version
HEAD^	One commit ago
HEAD^^	Two commits ago
HEAD~1	One commit ago
HEAD~3	Three commits ago
:/“Reformatting all”	Nearest commit whose comment begins with “Reformatting all”
RELEASE-1.0	Tag applied to the code when it was certified for release.



The complete set of revision specifications can be viewed by typing: **git help rev-parse**

Treeish can be used in combination with all Git commands that accept a specific commit or range of commits.

Examples include:

```
git log HEAD~3..HEAD
git checkout HEAD^^
git merge RELEASE-1.0
```

THE TYPICAL LOCAL WORKFLOW

Editing

Once you’ve cloned or initialized a new Git project, just start changing files as needed for your current assignment. There is no pessimistic locking of files by teammates. In fact, there’s no locking at all. Git operates in a very optimistic manner, confident that its merge capabilities are a match for any conflicted changes that you and your colleagues can craft.

If you need to move a file, Git can often detect your manual relocation of the file and will show it as a pending “move.” However, it is often more prudent to just directly tell Git to relocate a file and track its new destination.

```
git mv originalfile.txt newsubdir/newfilename.txt
```

If you wish to expunge a file from the current state of the branch, simply tell Git to remove it. It will be put in a pending deletion state and can be confirmed and completed by the next commit.

```
git rm fileyouwishtodelete.txt
```

Viewing

Daily work calls for strong support of viewing current and historical facts about your repository, often from different, perhaps even orthogonal points of view. Git satisfies those demands in spades.

Status

To check the current status of a project’s local directories and files (modified, new, deleted, or untracked) invoke the status command:

```
git status
```

```
[polyglot_maven_graven_git (master)]> git s
# On branch master
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
#       new file:   pom.csh
#
# Changed but not updated:
#   (use "git add/rm <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
#       modified:   pom.xml
#       deleted:    pom.yml
#
# Untracked files:
#   (use "git add <file>..." to include in what will be committed)
#
#       pom.bsh
[polyglot_maven_graven_git (master)]>
```

Diff

A patch-style view of the difference between the currently edited and committed files, or any two points in the past can easily be summoned. The .. operator signifies a range is being provided. An omitted second element in the range implies a destination of the current committed state, also known as HEAD:

```
git diff
git diff 32d4..
git diff --summary 32d4..
```

Git allows for diffing between the local files, the stage files, and the committed files with a great deal of precision.

Command	Definition
git diff	everything unstaged diffed to the last commit
git diff --cached	everything staged diffed to the last commit
git diff HEAD	everything unstaged and staged diffed to the last commit

```
[polyglot_maven_graven_git (master)]> git diff
diff --git a/pom.xml b/pom.xml
index 0b82ecc..9abeb36 100644
--- a/pom.xml
+++ b/pom.xml
@@ -93,7 +93,7 @@
<artifactId>apache-maven</artifactId>
<version>${mavenVersion}</version>
<classifier>bin</classifier>
- <type>zip</type>
+ <type>jar</type>
</dependency>

<dependency>
```

Log

The full list of changes since the beginning of time, or optionally, since a certain date is right at your fingertips, even when disconnected from all networks:

```
git log
git log --since=yesterday
git log --since=2weeks
```

```
[polyglot_maven_graven_git (master)]> git log --since=2weeks
commit 64de179becc3ed324daab72f7238df1404723672
Author: Jason Dillon <jason@planet57.com>
Date:   Tue Mar 16 14:07:04 2010 -0700

    Drop unused file
[polyglot_maven_graven_git (master)]>
```

Blame

If trying to discover why and when a certain line was added, cut to the chase and have Git annotate each line of a source file with the name and date it came into existence:

```
git blame <filename>
```

```
cbcb56e (Jason Dillon) 2009-11-27 14:55:36 +0700 26)      <version>5</version>
cbcb56e (Jason Dillon) 2009-11-27 14:55:36 +0700 27)      </parent>
f826e2eb (Mark Derricutt) 2009-12-06 00:14:49 +1300 28)
a4f22b9f6 (Jason Dillon) 2009-10-15 16:15:20 +0700 29)      <groupId>org.sonatype
a4f22b9f6 (Jason Dillon) 2009-10-15 16:15:20 +0700 30)      <artifactId>pmaven</a
0e5bfbd5 (Jason Dillon) 2009-10-14 01:40:49 +0700 31)      <name>Polyglot Maven<
3747fb57 (Jason Dillon) 2010-01-19 23:18:05 +0700 32)      <version>0.8-SNAPSHOT
a997212a (Heiko Seeberger) 2009-12-22 11:00:52 +0100 33)
0e5bfbd5 (Jason Dillon) 2009-10-14 01:40:49 +0700 34)      <packaging>pom</packa
fa2f40a5 (Jason Dillon) 2009-10-13 16:05:33 +0700 35)
10da6fdd (Jason Dillon) 2010-01-07 19:59:20 +0700 36)      <licenses>
10da6fdd (Jason Dillon) 2010-01-07 19:59:20 +0700 37)      <license>
10da6fdd (Jason Dillon) 2010-01-07 19:59:20 +0700 38)      <name>The Apa
10da6fdd (Jason Dillon) 2010-01-07 19:59:20 +0700 39)      <url>http://w
10da6fdd (Jason Dillon) 2010-01-07 19:59:20 +0700 40)      <distribution
10da6fdd (Jason Dillon) 2010-01-07 19:59:20 +0700 41)      </license>
10da6fdd (Jason Dillon) 2010-01-07 19:59:20 +0700 42)      </licenses>
cbcb56e (Jason Dillon) 2009-11-27 14:55:36 +0700 43)
```

Stashing

Git offers a useful feature for those times when your changes are in an incomplete state, you aren’t ready to commit them, and you need to temporarily return to the last committed (e.g. a fresh checkout). This feature is named “stash” and pushes all your uncommitted changes onto a stack.

```
git stash
```

When you are ready to write the stashed changes back into the working copies of the files, simply pop them back of the stack.

```
git stash pop
```

Aborting

If you want to abort your current uncommitted changes and restore the working copy to the last committed state, there are two commands that will help you accomplish this.

```
git reset --hard
```

Resetting with the hard option recursively discards all of your currently uncommitted (unstaged or staged) changes.

To target just one blob, use the checkout command to restore the file to its previous committed state.

```
git checkout -- Person.java
```

Adding (Staging)

When the developer is ready to put files into the next commit, they must be first staged with the add command. Users can

navigate to any directory, adding files item by item, or by wildcard.

```
git add <file name, folder name, or wildcard>
git add submodule1/PrimaryClass.java
git add .
git add *.java
```



Specifying a folder name as the target of a git add recursively stages files in any subdirectories.

The `-i` option activates interactive add mode, in which Git prompts for the files to be added or excluded from the next commit.

```
git add -i
```

The `-p` option is a shortcut for activation of the **patch** sub-mode of the interactive prompt, allowing for precise pieces within a file to be selected for staging.

```
git add -p
```

Committing

Once all desired blobs are staged, a commit command transactionally saves the pending additions to the local repository. The default text \$EDITOR will be opened for entry of the commit message.

```
git commit
```

To supply the commit message directly at the command prompt:

```
git commit -m"<your commit message>"
```

To view the statistics and facts about the last commit:

```
git show
```

If a mistake was made in the last commit's message, edit the text while leaving the changed files as-is with:

```
git amend
```

Branching

Branching superficially appears much the same as it does in other version control systems, but the difference lies in the fact that Git branches can be targeted to exist only locally, or be shared with (pushed to) the rest of the team. The concept of inexpensive local branches increases the frequency in which developers use branching, opening it up to use for quick private experiments that may be discarded if unsuccessful, or merged onto a well-known branch if successful.

```
git branch <new branch name> <from branch>
git branch <new branch name>
```

Choosing a Branch

Checking out (switching to) a branch is as simple as providing its name:

```
git checkout <branch name>
```

Local and remote git branches are checked out using the same command, but in somewhat of a radical change of operation for users coming from other systems like Subversion, remote branches are read-only until "tracked" and copied to a local branch. Local branches are where new work is performed and code is committed.

```
git branch <new branch name> <from branch>
git checkout <new branch name>
```

or alternatively, in a combined command:

```
git checkout -b <new branch name> <from branch>
```


Starting with Git 1.6.6, a shorthand notation can be used to track a remote branch with a local branch of exactly the same name when no local branch of that name already exists and only one remote location is configured.

```
git checkout <remote and local branch name>
git checkout performanceexperiment
```

Listing Branches

To list the complete set of current local and remote branches known to Git:

```
git branch -a
```



The local branches typically have simple names like master and experiment. Local branches are shown in white by Git's default syntax highlighting. Remote branches are prefixed by "remotes" and are shown in red.

Merging

Like other popular VCSes, Git allows you to merge one or more branches into the current branch.

```
git merge <branch one>
git merge <branch one> <branch two>
```

If any conflicts are encountered, which is rare with Git, a notification message is displayed and the files are internally marked with >>>>>>>> and <<<<<<<< around the conflicting portion of the file contents. Once manually resolved, git-add the resolved file, then commit in the usual manner.

Rebase

Rebasing is the rewinding of existing commits on a branch with the intent of moving the "branch start point" forward, then replaying the rewound commits. This allows developers to test their branch changes safely in isolation on their private branch just as if they were made on top of the mainline code, including any recent mainline bug fixes.

```
git rebase <source branch name>
git rebase <source branch name> <destination branch name>
```

Tagging

In Git, tagging operates in a simple manner that approximates other VCSes, but unlike Subversion, tags are immutable from a commit standpoint. To mark a point in your code timeline with a tag:

```
git tag <tag name>
git tag <tag name> <treeish>
```

THE REMOTE WORKFLOW

Working with remote repositories is one of the primary features of Git. You can push or pull, depending on your desired workflow with colleagues and based on the repository operating system file and protocol permissions. Git repositories are most typically shared via SSH, though a lightweight daemon is also provided.

Git repository sharing via the simple daemon is introduced at <http://www.kernel.org/pub/software/scm/git/docs/git-daemon.html>

Hot Tip

Sharing over SSH and Gitis is documented in the Git Community Book at http://book.git-scm.com/4_setting_up_a_private_repository.html

Remotes

While full paths to other repositories can be specified as a source or destination with the majority of Git commands, this quickly becomes unwieldy and a shorthand solution is called for. In Git-speak, these bookmarks of other repository locations are called **remotes**.

A remote called **origin** is automatically created if you cloned a remote repository. The full address of that remote can be viewed with:

```
git remote -v
```

To add a new remote name:

```
git remote add <remote name> <remote address>
git remote add <remote name> git@github.com:matthewmccullough/ts.git
```

Push

Pushing with Git is the sending local changes to a colleague or community repository with sufficiently open permissions as to allow you to write to it. If the colleague has the pushed-to branch currently checked out, they will have to re-checkout the branch to allow the merge engine to potentially weave your pushed changes into their pending changes.

```
git push
git push <remote name> <branch name>
git push <remote name> <local branch name:remote branch name>
```

```
[sample11-gitscm_git (master)]> git commit -a -m'Adding description'
[master 82abdf9] Adding description
1 files changed, 2 insertions(+), 1 deletions(-)
[sample11-gitscm_git (master)]> git push
Counting objects: 9, done.
Delta compression using up to 2 threads.
Compressing objects: 100% (6/6), done.
Writing objects: 100% (6/6), 785 bytes, done.
Total 6 (delta 4), reused 0 (delta 0)
To git@github.com:matthewmccullough/maven-training-samplegitscm.git
 f07c02b..82abdf9 master -> master
[sample11-gitscm_git (master)]>
```

Fetch

To retrieve remote changes without merging them into your local branches, simply **fetch** the blobs. This invisibly stores all retrieved objects locally in your .git directory at the top of your project structure, but waits for further explicit instructions for a source and destination of the merge.

```
git fetch <remote name>
git merge <remote name/remote branch>
```

Pull

Pulling is the combination of a fetch and a merge as per the previous section all in one seamless action.

```
git pull
git pull <remote name>
git pull <remote name> <branch name>
```

Bundle

Bundle prepares binary diffs for transport on a USB stick or via email. These binary diffs can be used to "catch up" a repository that is behind otherwise too stringent of firewalls to successfully be reached directly over the network by push or pull.

```
git bundle create catchupsusan.bundle HEAD~8..HEAD
git bundle create catchupsusan.bundle --since=10.days master
```

These diffs can be treated just like any other remote, even though they are a local file on disk. The contents of the bundle can be inspected with ls-remote and the contents pulled into the local repository with fetch. Many Git users add a file extension of .bundle as a matter of convention.

```
git ls-remote catchupsusan.bundle
git fetch catchupsusan.bundle
```


GUIs

Many graphical user interfaces have gained Git support in the last two years. The most popular Ruby, Perl, and Java/JVM IDEs have between a good and great level of Git integration today.

Gitk & Git Gui

Standard Git distributions provide two user interfaces written in Tcl/Tk. The first, Git-Gui offers a panel by which to select files to add and commit, as well as type a commit message. The latter offers a diagram visualization of the project's code history and branching. They both assume the current working directory as the repository you wish to inspect.

```
git gui
gitk
```



IDEs

Java IDEs including IntelliJ, Eclipse (eGit), and NetBeans (NBGit) all offer native or simple plugin support for Git through their traditional source code control integration points.

Numerous other platform-native GUIs offer graphically rich history browsing, branch visualization, merging, staging and commit features.

Hot Tip

A complete list of Git IDEs and GUIs can be found at: <http://delicious.com/matthewmccullough/git+gui>

CVS, SUBVERSION

On the interoperability front, the most amazing thing about Git is its ability to read and write to a remote Subversion or CVS repository while aiming to provide the majority of the benefits of Git on the local copy of a repository.

Cloning

To convert a Subversion repository that uses the traditional trunk, tags, branches structure into a Git repository, use a syntax very similar to that used with a traditional Git repository.

```
git svn clone --stdlayout <svn repo url>
```

Please be patient, and note the progress messages. Clones of large Subversion repositories can take hours to complete.

Pushing Git Commits to Subversion

Git commits can be pushed, transactionally, one for one to the cloned Subversion repository. When the Git commits are a good point for sharing with the Subversion colleagues, type:

```
git svn dcommit
```

Retrieving Subversion Changes

When changes are inevitably made in Subversion and it is desired to freshen the Git repo with those changes, rebase to the latest state of the Subversion repo.

```
git svn rebase
```

ADVANCED COMMANDS

Git offers commands for both the new user and the expert alike. Some of the Git features requiring in-depth explanations can be discovered through the resources links below. These advanced features include the embedded (manpage-like) **help** and ASCII art visualization of branch merge statuses with **show-branch**. Git is also able to undo the last commit with the **revert** command, binary search for (**bisect**) the commit over a range of history that caused the unit tests to begin failing, check the integrity of the repository with **fsck**, prune any orphaned blobs from the tree with **gc**, and search through history with **grep**. And that is literally just the beginning.

This quick overview demonstrates what a rich and deep DVCS Git truly is, while still being approachable for the newcomer to this bold new collaborative approach to source code and version control.

REFERENCES

Git has a rich set of tutorials, screencasts and published books to satisfy your newfound thirst for knowledge on all things related to this unique DVCS.

Official Homepage, Releases

For official release notes and 'man' pages, take a look at:

<http://git-scm.com/> or <http://www.kernel.org/pub/software/scm/git/docs/>

Manuals, Tutorials

With the growing interest around Git, well-crafted tutorials and how-to guides are sprouting up across the 'net. Including:

<http://cworth.org/hgbook-git/tour/>

<http://www-cs-students.stanford.edu/~blynn/gitmagic/>

<http://www.kernel.org/pub/software/scm/git/docs/user-manual.html>

<http://peepcode.com/products/git>

Books

Pro Git by Scott Chacon. Free in HTML form. Also available in printed form from Apress. <http://progit.org/book/>

Version Control with Git by Jon Loeliger. Printed by O'Reilly.

<http://oreilly.com/catalog/9780596520137>

Pragmatic Version Control Using Git by Travis Swicegood.

Printed by Pragmatic Programmers.

<http://pragprog.com/titles/tsgit/pragmatic-version-control-using-git>

Bookmarks

Given the rapid pace of changes to the Git ecosystem, I maintain a hand-culled list of the most informative pages on the Internet for the benefit of the Git community.

<http://delicious.com/matthew.mccullough/git>


ABOUT THE AUTHOR



Matthew McCullough is an Open Source Architect with the Denver, Colorado consulting firm Ambient Ideas, LLC which he co-founded in 1997. He's spent the last 13 years passionately aiming for ever-greater efficiencies in software development, all while exploring how to share these practices with his clients and their team members. Matthew is a nationally touring speaker on all things open source and has provided long term mentoring and architecture services to over 40 companies

ranging from startups to Fortune 500 firms. Feedback and questions are always welcomed at matthewm@ambientideas.com


RECOMMENDED BOOKS



BUY NOW



BUY NOW



BUY NOW



Browse our collection of over 100 Free Cheat Sheets

Free PDF

Upcoming Refcardz

OSMF

Clojure

HTML 5

Test Driven Development



DZone communities deliver over 6 million pages each month to more than 3.3 million software developers, architects and decision makers. DZone offers something for everyone, including news, tutorials, cheatsheets, blogs, feature articles, source code and more.

"DZone is a developer's dream," says PC Magazine.

DZone, Inc.

140 Preston Executive Dr.

Suite 100

Cary, NC 27513

888.678.0399

919.678.0300

Refcardz Feedback Welcome

refcardz@dzone.com

Sponsorship Opportunities

sales@dzone.com



\$7.95